

THE FUNDAMENTAL TUG-OF-WAR IN PHYSICAL AI

Reconciling High-Capacity Stochastic Deliberation with Deterministic Real-Time Physics

AI DELIBERATION REALM

Cognitive Capacity

Demands Compute, Tokens & Time

Computational Nature:

- High-capacity VLMs & Diffusion policies
- Billions of parameters & KV cache matrices
- Non-deterministic stochastic outputs

Temporal Profile:

Slow Cadence: 0.5–50 Hz

P99 Latency Tails $\approx$ 500–1500 ms

Operating substrate: Linux MPU / Edge NPU.

STRUCTURAL TENSION

Latent Time vs. Physical Reality

PHYSICAL DYNAMICS REALM

Kinetic Determinism

Demands Microseconds & Hard Bounds

Physical Reality:

- Kinetic momentum $p = mv$ & inertia
- Dynamic stopping distance $d_{\text{stop}}(v)$
- Actuator Joule heating & gear shear

Temporal Constraint:

Fast Reflex: 1000 Hz (1.0 ms)

Hard Deadline: $\tau_{\text{world}} \leq 1\text{--}10\text{ ms}$

Operating substrate: Real-Time Bare-Metal MCU.

HOW PHYSICAL AI RECONCILES THE TUG-OF-WAR: THREE ARCHITECTURAL CONTRACTS

1. Proposal–Permission Privilege Split

- Untrusted MPU proposals never touch motors.
- Dedicated 1 kHz MCU filters commands via Control Barrier Functions: $h(x) \geq 0$.

2. Action Chunking (System 1.5)

- Amortizes neural latency over future horizon H .
- Synthesizes continuous C^2 jerk-bounded trajectories executed locally at 1 kHz.

3. Expiring Intent Leases (L_{intent})

- Binds proposals to time-to-live validity (TTL).
- Continuous state tracker revokes lease on drift, commanding deterministic safe position hold.